

Team Project

2026 Spring, SWPP

All scheduled dates are tentative.

Goal

In this project, you'll work as a team with other students to write a simple compiler that reads an LLVM IR program and emits an optimized assembly for an imaginary machine.

Each team should

- Use LLVM **22.1.0** built with the script in class repo for your project.
- Use GitHub to collaborate and code-review teammates' implementation.

At the end of the semester, we'll have a competition to evaluate the performance of each team's compiler.

Updates

- Apr. 10 (Fri) : Change the document submission method, explicitly specify the team repository name, clarify policy about AI collaborator's review
- Apr. 28 (Tue) : Add explicit notice about modifying the code out of scope, instruct each team member to create their own fork at the beginning of the project, add detailed policy about 'syntax-based optimization', fix incorrect text
- May. 4 (Mon) : Add explicit notice about postponing the PR to the later sprint, add rule for unit test exemption in limited cases, add rule for adaptor pass for existing LLVM pass
- May. 6 (Wed) : Relax page limits for progress report and reqspec, reduce minimum review period to 48 hours
- May. 6 (Wed) : Change the typo about the due date in the schedule section.
- May. 12 (Tue) : Modify the Team score policy into *Collaboration* score.

Schedule (subject to change)

Date (Tentative)	Details
4/16	Introduce your team in class
4/17 - 4/23	Sprint 0 - Read codes, setup development envs and CI, etc.
4/23	Submit two documents: - Requirement and specification - Planning
4/24 - 5/7	Sprint 1 - ~10 days (From Friday): development - 3+ days (until Wednesday): code review+revision - Last day: progress report & req/spec plans
5/8 - 5/21	Sprint 2
5/22 - 6/4	Sprint 3
6/5 - 6/8	Wrap-up
6/9	Competition

(1) Sprint 0

Introduce your team in the class

Each team will have about 5 minutes to introduce themselves in the class.

- Introduce your team & teammates.
- Explain what you want to learn during the team project about software engineering.
- Suggest 5+ optimizations (either simple or complex) with example codes.
- Prepare short presentation slides.
- The talk is lightweight; your talk does not need to get very technical.

Pre-project Documentation

You should upload two documents (A. requirement and specification, B. planning) before the first sprint. You should upload the document on a **team repository issue board**. Any team member can upload the document on behalf of their team. Each document should not be more than 7 pages. You may use Korean or English for these documents. The document should be in **pdf** format.

A. Planning

- The title for the issue should be “[Sprint 0, Document] Planning”
- Upload your document as a pdf file. Don’t write it in the article!

- If we notice that the issue has been edited after the submission deadline, we will mark the submission as zero points.
- File name should be “sprint0-teamNN-planning.pdf” (N is your team number)
- Briefly explain optimizations that your team would like to implement during the 3 sprints.
- It should include each member’s plan for 3 sprints.
- Optimizations can be implemented with your teammates.
- Optimizations can be implemented through multiple sprints.

B. Requirement and specification

- The title for the issue should be “[Sprint 1, Document] Requirements and Specifications”
- Upload your document as a pdf file. Don’t write it in the article!
- If we notice that the issue has been edited after the submission deadline, we will mark the submission as zero points.
- File name should be “sprint1-teamNN-reqspec.pdf” (note the sprint number)
- Describe the optimizations that you are going to implement in sprint 1. For each optimization, describe the following information.
 - i. Description
 - ii. An algorithm in a high-level pseudo-code
 - iii. At least 3 pairs of IR programs, before and after the suggested optimization

Preparing the Project Skeleton

Once the team project repository template has been released, each team should create their own team repository from the template provided. The repository must be named **swpp202601-teamNN** (e.g. team01, team22, ...). During sprint 0, students can directly push commits that contain project skeleton code, CMakeLists configuration, or GitHub Actions configuration into the team repository. These commits are exempt from PR criteria (described below) and grading.

Especially, automated testing on the GitHub Actions should be fully functional by the beginning of sprint 1. Details about testing are described later in this document.

After the team repository has been created, each team member, other than the team repository owner, should create an individual fork from the team repository, and work on their respective repository.

(2) Sprint 1/2/3

Each sprint consists of

- **Development phase** (Up to 11 days, starting from Friday)
 - Implement the planned features
 - Send pull requests to the main repository
 - Assign reviewers.
- **Code-review phase** (At least 48 hours, until Wednesday)

- Assigned reviewers review the pull requests
- Reviewee updates the pull request according to the comments.
- If the update is complete, the pull request is merged into the main repository
- **Documentation phase** (On Thursday)
 - Write a progress report for this sprint, and requirements and specifications for the next sprint. Upload them as team repository Issues.

You may use Korean or English in each process.

Pull Requests

- A pull request contains an implementation of a single feature.
- The title of a pull request should start with [Sprint N]. (N is the current sprint number)
- It should contain unit tests to check the added functionality.
- Its line diff should be at most around 500 lines except comments, spaces, tests, and non-C++ code files such as .gitignore or CMakeLists.txt.
 - Use a formatter whenever you write code. It's also helpful to use Git hooks or similar tools to automatically format your code before submitting a PR. You may use any formatter, but we recommend using the one we provided in the .clang-format file.
 - We'll count the LOC after running the formatter on your code. Formatting rules can be found in included .clang-format
- It should satisfy the good pull request conditions that are described in Apr.9th's practice session
- You may merge to any branch on the repository. However, only the PRs that are merged into the **main** branch will be graded.
 - If the opt. pass PR is to be merged to the main branch, it must be reviewed.
 - Otherwise, reviews are optional but recommended.
- Each pull request should be merged using 'squash and merge' by the pull request writer (not reviewers!)
- Detailed information about PR grading is provided at the end of this document.

Policy for Writing Pull Requests & Reviews

- Each student can merge up to 6 pull requests per sprint.
 - If your team has less than four teammates for any reason, you may take the 'merge chance' of the missing teammate(s) as well.
- Each student should merge at least one pull request per sprint.
 - Only 'graded' PRs count; otherwise, we cannot grade your performance on a sprint and have to give you zero points.
 - We will count co-authored PR as one PR per each person.
- For each pull request, 2 or more human reviewers should be assigned.
- You may set up automated AI collaborator's review, but it won't count as 2 reviews requirement.

- Every PR must pass the CI at the time of merge.
 - Merging the PR that fails the CI (either on PR or after merge) will be penalized, with no exception.
 - Each commit onto 1) pull request (on: pull_request) to the team repo, and 2) main branch of the team repo should trigger GitHub Actions and run the CI.
 - When your PR is based on 'older' main, your PR may pass the CI but it may fail upon merging onto the main branch. Make sure to merge the content of the main branch and re-run the CI before merge.
- If necessary, a team can write pull requests that make **non-functional changes** such as directory structure changes / source file splitting / etc.
 - Please add "[Sprint N, NFC]" at the beginning of the title.
 - This pull request should be reviewed but not as much as functional ones.
 - This pull request is exempt from the 6 pull request restriction.
 - We'll be generous about what counts as 'non-functional':
As long as your PR does not change the function of the compiler code, you may call it NFC.
- If necessary, a team can make reverting pull requests that **revert previous PR**
 - Please add "[Sprint N, Revert]" at the beginning of the title.
 - Please add "[Sprint N, Reverted]" at the title of the reverted PR.
 - This pull request should be reviewed but not as much as functional ones.
 - 'Revert' and 'Reverted' pull requests are exempt from the 6 pull request restriction.
 - While 'Revert' PR will not be graded, we will still use 'Reverted' PR for grading.
- **TAs' codebase update** should be applied in the form of PR as well.
 - Please add "[Sprint N, Update]" at the beginning of the title.
 - This pull request does not have to be reviewed, but it should still pass the CI!
 - You may apply Update PRs at any moment of the project period, but all updates must be applied to the team repository by the end of wrap-up period, as they may contain critical bug fixes.
- Other than TAs' codebase update, the *content* of source code that is out of project scope must not be modified.
 - Our interpretation of *source code* is 'files that get involved when compiling the swpp-compiler binary'.
 - You are allowed to add, modify or remove any file or directory inside the project root directory, except src, test, and CMakeLists.txt.
 - You are not allowed to add, edit, or remove any file or directory inside src, except opt.cpp and opt directory.
 - You are allowed to add, modify or remove any file or directory inside test directory, including CMakeLists.txt.
 - Your repository should be able to build swpp-compiler and run ctest
 - Such an attempt will be penalized by zero points on your **entire** project score.
 - We will compare the content with the skeleton code at the end of the project, and apply the penalty if we detect any difference other than formatting.

- The PR doesn't have to be merged in the very sprint it has been initially created.
 - In case a PR has been created in a sprint and is planned to be reviewed & merged in the later sprint, the title tag ([Sprint N]) must be changed to the sprint that the PR is planned to be merged.

Existing Optimizations

LLVM has many optimization passes, and simply calling these functions from your project can sometimes bring large performance improvements.

Each existing optimization should be added via separate PR, and should come with unit tests as well. This does count in six PR limits. Adding an existing pass as NFC will be penalized.

If and only if the existing LLVM pass yields an IR program that cannot be accepted by the compiler backend, you may append an 'adaptor pass' into the PR. The function of adaptor pass should be limited to code cleanup, and the core optimization should only depend on the existing pass.

Existing Analyses

LLVM has many analysis passes as well.

Unlike optimization passes (that actually modify the code structure), you may include and use the analysis passes as part of your pass implementation. Only the pass that you have implemented will count in six PR limits.

Writing & Running Unit Tests for Optimization Pass

- Each optimization pass PR should contain at least 3 unit tests.
- Unit tests are responsible for checking the functionality of 'single' optimization pass.
 - Testing the output of `opt --<new-opt-pass>` is therefore recommended
- Each unit test should be an LLVM IR program with FileCheck comments.
- Running `ctest --test-dir build` should run all added unit tests.
- Adding Alive2 tests is not mandatory but is highly recommended.
- When adding the same pass into different stages of the optimization pipeline, without changing the pass itself, unit tests may be omitted and you may put a link to the initial PR that contains the unit test instead.

Writing & Running Unit Tests for Analysis Pass

- Each analysis pass PR should contain at least 3 unit tests.
- Unlike optimization passes, analysis results are not visible in IR programs, and requires direct access into LLVM's internal data structure to check its functionality.
- Therefore, each unit test should be a C++ code that loads an LLVM IR program, runs the program via analysis pass, and asserts that your pass functions as expected.
- Other than the implementation method, the rule for optimization pass unit tests apply the same.

Main Repository

- The name of the main repository must be **swpp202601-teamNN**
- The main repository should contain a **main** branch.

- After each commit, the project should work correctly. TA will check
`cmake --build build --target swpp-compiler # we will modify related variables as necessary`
`ctest --test-dir build`
 from the root directory of the branch.

Using GitHub Issue

- Please use the GitHub issue to show that you are actively communicating with people.
- If there was an offline discussion and it wasn't written as comments at Pull Request, please record it at GitHub Issue.

Program-Specific Optimization

- Optimizing a program based on function name matching is not allowed.
- Replacing an entire program implementation with pre-written code is not allowed.
- Your optimization should not depend on specific syntax that only appears in a limited set of programs.
 - 'Syntax' means name of function, basic block, and virtual register.
 - That is, your optimization should be able to emit equivalent code even after we obfuscate the program.
- We will obfuscate the program before running the compiler, by renaming functions, basic blocks, and virtual registers.
 - But we will make changes only to the extent that do not break the logical structure, core assumptions about valid SWPP program, and LLVM grammar.
 - For example we won't rename main function, entry block, read() and write(i64) calls, etc.
- Even if your optimization fail to satisfy the above conditions, we will not give penalty for your implementation. It just won't improve the performance as expected.
 - But if your optimization's core idea is about exploiting the syntactic hint from the original program, we will deduct the explanation score from your PR.
 - Note that we only restrict using hint from the program before any optimization. Using the syntax of emitted IR from other passes as hint can be useful in some cases.

End-of-the-Sprint Documentation

At the end of each sprint, each team should submit the progress report and the updated requirements/specifications for the next sprint. The document should be in **pdf** format.

A. Progress report

- The title for the issue should be “[Sprint N, Document] Progress” (N is the **current** sprint)
- Upload your document as a pdf file. Don't write it in the article!
- If we notice that the issue has been edited after the submission deadline, we will mark the submission as zero points.
- File name should be “sprintN-teamMM-progress.pdf”

- Describe each member's progress at the end of the sprint.
- If you did not finish what you have planned, explain why.
- Include the results of applying your passes to all benchmarks (distributed at the beginning of the project) as well as your own tests to show the effectiveness of your optimizations. Examples may include performance enhancements, applicability (how many cases do your passes have effects), etc.
- If there is an update in planning, please submit the updated part as well.
- This document should not exceed 8 pages.

B. Requirement and specification

- The title for the issue should be "[Sprint N, Document] Requirements and Specifications" (N is the **upcoming** sprint)
- Upload your document as a pdf file. Don't write it in the article!
- If we notice that the issue has been edited after the submission deadline, we will mark the submission as zero points.
- File name should be "sprintN-teamMM-reqspec.pdf"
- Describe optimizations that you are going to implement in the next sprint. For each optimization, description / pseudo-code / IR programs should be included.
- This document should not exceed 10 pages.

You should submit two documents as team repository issues. You may use Korean or English.

	~ Sprint 0	~ Sprint 1	~ Sprint 2	~ Sprint 3
	4/23 11:59 pm	5/7 11:59 pm	5/21 11:59 pm	6/4 11:59 pm
Planning	O			
Requirement and specification	O (1)	O (2)	O (3)	
Progress report		O (1)	O (2)	O (3)

(3) Wrap-up & Competition

After 3 sprints, you will be given a few days to wrap up your implementation.

In this period, you may commit codes as much as you want (no code review, no line diff constraint, ...)

On the last day, we'll run a competition, estimating the correctness & efficiency of the compiler.

AI Policy

The use of AI is allowed for code, pull requests, review comments, and project documents. However, students are fully responsible for the correctness and quality of anything they submit. Before submission, you must carefully check whether the AI-generated content is accurate, technically sound, and appropriate for the project. Submissions that contain incorrect, misleading, or unnatural AI-generated content may receive a substantial penalty.

If you used AI in a pull request, review, or document, this must be explicitly disclosed.

In particular, students must not modify code based only on function names or superficial pattern matching. Code should be changed only when the student actually understands its behavior and purpose.

Sprint PR Grading Policy

This project has three sprints. For each sprint, the score is composed of:

- **Personal score: 20 points**
- **Collaboration score: 5 points**

So, each sprint is worth 25 points.

Personal Score (20 points)

The personal score for a sprint is based on the average score of that student's PRs in the sprint.

- **Minimum PR Requirement**
 - A student must submit at least 2 PRs in the sprint to be eligible for the full 20-point individual score.
 - If a student submits only 1 PR, the individual score is capped at 10 points, even if the PR itself receives a higher score.
 - If a student submits 0 PRs, the individual score is 0 points.

Grading PR

Each PR is graded using the following rubric:

- **Explanation: 10 points**
- **Performance Evaluation: 4 points**
- **Correctness: 8+1 points**

The raw maximum for a PR is therefore 23 points, but the recorded PR score is capped at 20 points. In other words, Recorded PR score = $\min(\text{Raw PR score}, 20)$.

PR Rubric

1) Explanation (10 points)

This category evaluates whether the PR clearly explains the problem and the implementation.

a) Issue Description

- The PR should briefly explain what the current issue is.
- **Full credit:** The PR clearly states what problem, inefficiency, or missed optimization opportunity exists.
- **No credit:** The PR does not describe the issue at all, or only says that something was changed without explaining why.

b) Implementation / High-Level Algorithm

- The PR should explain how the pass is implemented.
- **Full credit:** The PR describes the main implementation idea clearly. It includes a high-level algorithm description or equivalent explanation of the core transformation logic.
- **Partial credit:** The PR explains the implementation only partially, or at a very superficial level.

- **Low credit:** The reviewer must infer the implementation only from the diff.

2) Performance Evaluation (4 points)

This category evaluates whether the PR includes quantitative evidence and analysis.

The amount of speedup itself is **not** part of the grading criteria. What matters is whether the PR includes **quantitative evidence** and a **reasonable analysis**.

a) Quantitative Result

- The PR should include numerical performance results.
- **Full credit:** The PR clearly states what benchmark(s) or cases were tested and includes a table, figure, or clearly written numerical comparison.
- **No credit:** The PR only claims that performance improved, with no numbers.

b) Analysis

- The PR should include some interpretation of the result.
- **Full credit:** The PR explains what the result means, even briefly.
- **No credit:** Numbers are given with no discussion.

3) Correctness (8+1 points)

This category evaluates whether the PR provides evidence that the optimization is correct.

a) Unit Tests and Example Program Tests

- The PR must show that **both** unit tests and example program tests were run successfully.
- The PR should explain what was tested.
- **Full credit:** The PR clearly shows that both unit tests and example program tests passed. The PR states which tests were run, or gives a sufficiently clear summary of test coverage.
- **Partial credit:** The PR says tests were run, but the scope is unclear.
- **No credit:** No meaningful test report is provided.
- **Important rule:** If either one is missing, this item receives 0 points.

b) Correctness Discussion

- The PR should provide additional evidence that the optimization is safe.
- Examples of acceptable evidence:
 - discussion of safety conditions,
 - reasoning about edge cases,
 - explanation of why the transformation preserves correctness,
 - limitations intentionally left unoptimized.
- **Full credit:** The PR gives a technically meaningful correctness discussion.

- **Partial credit:** The PR includes only a weak or incomplete correctness argument.
 - **No credit:** No correctness discussion is provided.
- c) **Alive2 Bonus (Optional)**
- Alive2 is **not required**, but it can earn a bonus point.
 - **Full credit:** The PR reports meaningful Alive2 usage and result(s).
 - **No credit:** Alive2 is not used, or is mentioned without any meaningful evidence.

Collaboration Score (5 points)

Each student starts with **5 points per sprint** for the Collaboration Score. Points are deducted for process issues that negatively affect team collaboration. The Collaboration Score cannot go below **0 points** and cannot exceed **5 points**.

Deductions

Points may be deducted as follows:

- **CI did not pass when merging a PR:** -1 per incident
- **PR was not squash-merged:** -1 per incident
- **Direct push to the main branch:** -1 per incident
- **Merging another teammate's PR:** -1 per incident
 - In principle, each student should merge their own PR unless there is a clearly justified reason.
- **Merging a PR before the minimum review period of 48 hours has passed:** -1 per incident
 - A PR should normally remain open for at least 48 hours before being merged, so that teammates have enough time to review it.
 - If such PR was posted later than 48 hours before sprint deadline, **additional 1 point will be deducted**
- **Other behavior that does not contribute to the team project process:** additional deductions may be applied later
 - For example, consistently providing too few or insufficient reviews on teammates' PRs, not participating in the review process, or otherwise failing to support the team's collaboration may result in deductions.
 - A contribution survey may be conducted to assess each student's contribution to the team process.

Bonus

Students who contribute especially actively to the team process may receive bonus points, for example by providing many helpful reviews, supporting teammates, or improving the team's workflow.

However, the Collaboration Score is capped at **5 points**, so bonus points cannot raise the score above the maximum.

What counts as an insufficient PR review?

- A PR review is considered insufficient if it does not provide meaningful technical feedback.
- Examples of insufficient review:
 - “LGTM”
 - “Looks good”
 - “Good job”
 - emoji-only or approval-only comments
 - comments only about title, formatting, or trivial wording
- Examples of a sufficient review:
 - raising a correctness concern
 - asking about an edge case
 - questioning performance evaluation
 - asking whether tests are adequate
 - discussing algorithm design or implementation choices

General Notes

- A technically strong PR can still lose points if it is poorly explained.
- A well-written PR can still lose points if it lacks quantitative evidence or correctness validation.
- Claims without evidence will not receive full credit.
- The grading focuses on whether the PR communicates the work clearly and supports its claims with evidence.

Example PR

[Sprint 2] Strength reduction for multiplication by powers of two

Summary

This PR implements a simple strength-reduction optimization in our LLVM pass.

When the pass finds an integer multiplication by a constant power of two, it replaces the multiplication with a left shift by $\log_2(\text{constant})$. For example, $x * 8$ is transformed into $x \ll 3$ when the transformation is valid.

Issue Description

In the current implementation, expressions such as $x * 2$, $x * 4$, and $x * 8$ remain as integer multiplication instructions throughout our optimization pipeline. This is a missed optimization opportunity because multiplication by a power-of-two constant can often be expressed more simply as a shift.

Even if later backend stages may sometimes perform a similar simplification, it is still useful to canonicalize this pattern earlier in our pass pipeline. Doing so can reduce instruction cost, simplify IR, and expose additional opportunities for later optimizations.

Implementation / High-Level Algorithm

Main idea

The pass matches `mul` instructions where one operand is an integer constant and the other is an arbitrary integer value. If the constant is a positive power of two, the multiplication is replaced with an equivalent left shift.

High-level algorithm

1. Iterate through all instructions in each function.
2. Identify integer `mul` instructions.
3. Check whether one operand is a constant integer.
4. Check whether the constant is a positive power of two.
5. Compute the shift amount $k = \log_2(C)$.
6. Replace:
 - `mul x, C` with `shl x, k`
 - `mul C, x` with `shl x, k`
7. Leave all other cases unchanged.

Important implementation details

- The transformation is applied only to integer scalar types in this version.
- The pass only transforms cases where the constant operand is a positive power of two.
- Cases such as $x * 0$, $x * 1$, negative constants, non-power-of-two constants, and vector types are intentionally left unchanged in this PR.
- This keeps the transformation simple and avoids mixing this PR with other simplification rules.

Why this implementation is reasonable

A shift is a natural representation for multiplication by powers of two. Restricting the optimization to simple scalar integer cases makes the pass easier to validate and reduces the risk of introducing bugs in edge cases.

Performance Evaluation

Benchmarks

I evaluated the optimization on the provided example programs and several microbenchmarks containing repeated arithmetic patterns.

Quantitative Results

Benchmark	Baseline (ms)	Optimized (ms)	Change
arith_loop	121.4	113.6	6.4% faster
image_kernel	204.7	197.1	3.7% faster
matrix_scalar	88.9	85.8	3.5% faster
vector_ops	97.8	97.5	0.3% faster
simple_micro_mul_pow2	45.2	39.4	12.8% faster

Analysis

The optimization provides the clearest benefit on programs that contain many integer multiplications by compile-time powers of two that survive until our pass runs. This is why the improvement is most visible in `arith_loop` and `simple_micro_mul_pow2`.

The effect is smaller on programs such as `vector_ops`, where this pattern appears less often or is not on the hot path. In some cases, backend code generation may already perform a similar simplification, so the speedup is not always large. That is expected.

The goal of this PR is not to guarantee a large improvement on every benchmark, but to show that the optimization is quantitatively measurable and behaves as intended.

Correctness

Unit Tests

I added unit tests for the following cases:

- `x * 2 -> x << 1`
- `x * 4 -> x << 2`
- `x * 8 -> x << 3`
- `2 * x -> x << 1`
- non-power-of-two constants are not transformed
- `x * 0` is not handled by this PR
- `x * 1` is not handled by this PR
- negative constants are not transformed

- non-integer or unsupported cases remain unchanged

All added unit tests passed.

Example Program Tests

I ran the full example-program test suite after applying this optimization.

Test result summary:

- all provided example programs compiled successfully
- all example programs produced the same output as the baseline compiler
- no regression was observed in the existing example tests

So both the unit tests and the example program tests passed successfully.

Correctness Discussion

This transformation is safe in the restricted cases handled by this PR because multiplication by a positive power-of-two constant is semantically equivalent to a left shift by the corresponding amount for the same integer value representation used in this pass.

To keep the transformation conservative, this PR intentionally avoids broader cases. In particular:

- non-power-of-two constants are ignored
- negative constants are ignored
- unsupported types are ignored
- vector cases are left for future work

By restricting the optimization scope, the pass only rewrites cases that are easy to reason about and test.

Alive2

I also checked representative cases with Alive2.

Alive2 results:

- `mul i32 %x, 2 <-> shl i32 %x, 1`: verified
- `mul i32 %x, 8 <-> shl i32 %x, 3`: verified
- `mul i64 %x, 16 <-> shl i64 %x, 4`: verified

Alive2 was used only for representative scalar cases in this PR, but the results are consistent with the intended transformation rule.